


The Best Report Yet, and It Cites Itself as Evidence

Incident **15369769f5175238**, generated 24 Aug 2026. New tenant, new principal, new technique instance. Scored on the same fourteen metrics and weighting as the previous two.

56.
2

+8.6

Top of UNRELIABLE, four points below SUPERVISED. Eleven of fourteen metrics improved again. Response automation has been properly solved. The framework-to-substance ratio is now **1.15**, down from 2.04 at the first assessment -- the presentation layer and the evidence layer have essentially converged. This is a materially better product than three weeks ago.



THE THREE FINDINGS THAT MATTER

01 The corroborating "prior verdict" is this incident.

H1's fifth supporting item cites KB_RECORD_ID 15369769f5175238 as a prior high-confidence verdict for the same principal. **That is this report's own incident ID** -- printed at the top of page 1. The system retrieved its own record, described it as prior corroborating evidence, and used it to pass the false-positive check that exists specifically to detect uncorroborated conclusions. This was flagged as a projected risk in the last assessment. It is now demonstrated in print.

02 Response automation is genuinely, properly fixed.

Three actions -- Notify SOC, Create Ticket, Revoke Session -- each with a stated threshold, an explicit rationale, and human-approval gating. The Revoke Session note even shows its arithmetic: *"Tier 3 investigation raised confidence to 0.90, clearing this action's 0.75 threshold that the original 0.35 match confidence did not."* That is a real confidence-to-action policy matrix. It is the single best fix across all three releases and it closes the defect that sat at 12/100 in the first assessment.

03 T1685.002 is still fabricated. Fourth document, fourth appearance.

It has now migrated into the narrative prose on page 1 ("A defense evasion technique (T1685.002) was detected"). The correct mapping remains **T1562.008**. Credit where due: the invented technique *name* from the last build is gone, and Detection Opportunity 3 correctly cites T1078.004, a real ID -- so the system can handle valid identifiers. It simply has never been made to validate this one.

01 What this release fixed

Nine substantive fixes, several of them addressing defects that had survived two prior assessments. This is the strongest remediation pass of the series.

#	FIX	EVIDENCE IN THE REPORT	CLOSES
1	Recommended Actions are populated and threshold-gated.	Notify SOC, Create Ticket, Revoke Session, each with rationale and human approval required. "No actions were warranted" is gone.	D-06, N-02
2	Trail lifecycle state is now established.	get_trail_lifecycle_status returns currently_disabled and delete_trail_occurred: False. Two questions unanswered across two prior reports, both closed by one tool.	D-07, D-08
3	The source IP is actually enriched.	lookup_ip identifies 185.220.101.47 as a current_tor_exit_node. Both prior reports recommended this and neither did it.	--
4	Concrete artifacts appear.	Trail name northwind-eng-trail, region eu-west-1, a real routable IP. Prior reports named neither trail nor region.	D-14 (part)
5	Three-valued evidence logic appears for the first time.	FP check 3 reads "Not applicable" with a stated reason: the context-store check requires an asset_id not bound in this match. Previously this would have been counted as disconfirming evidence.	D-16 (part)
6	The Disposition is honest about being a single alert.	"This technique carries high enough potential impact to warrant immediate investigation on its own, without waiting for a second, corroborating alert." The v1 report falsely claimed corroboration here.	D-03
7	No fabricated chain.	All four correlation legs read "Proposed, not found," including discovery. The system did not invent a kill chain where none existed.	D-04
8	The invented technique name is gone.	The bogus "Inhibit System Recovery" label from the last build does not reappear.	N-01
9	A coverage gap is reported as a product gap.	Lessons Learned names the missing asset_id binding as "a gap in contextual data integration for cloud-native incidents" rather than treating the null as a finding.	--

The verdict is, for the first time in this series, actually supported by its evidence.

A Tor exit node performing StopLogging on a trail confirmed still disabled, with no replacement created, is a genuinely strong malicious case. The v1 report reached a probably-correct conclusion through absence-of-evidence reasoning. This one reaches a probably-correct conclusion through *evidence*. That is the distinction the first assessment said mattered most, and it has been made.

02 The self-citation

The previous assessment identified circular corroboration as a projected risk and recommended verdict provenance classes as a P0 fix. That fix was not made, and the risk has now materialised in the most literal form available.

WHERE	WHAT IT SAYS
Page 1, header	Incident 15369769f5175238
Page 2, H1 supporting evidence	"The get_identity_verdict_history shows a prior verdict (KB_RECORD_ID 15369769f5175238) for the same principal with a high investigation confidence (0.9) concluding defense evasion by a compromised principal from a Tor exit node. This strongly corroborates the current hypothesis. "
Page 3, FP check 4	Pass. "This alert is corroborated by two independent signals: the source IP [...] and the get_identity_verdict_history revealing a prior high-confidence malicious verdict for the same principal."
Page 4, Recommended Actions	"A principal whose own audit-trail-disabling action is corroborated by independent Tier 3 evidence (not the bare match alone) is the clearest real signal an active AWS session should be cut."

The identifiers are identical. The "prior verdict" and the current incident are the same record.

Its described content -- defense evasion, compromised principal, Tor exit node, confidence 0.9 -- is a word-for-word description of *this investigation's own conclusion at 0.90*. It is not a similar earlier case. It is a mirror. And it is not decorative: it is one of two signals that clear FP check 4, and FP check 4 is cited in the Recommended Actions block as the justification for revoking a live production session.

WHAT THIS ACTUALLY COSTS

- **The corroboration count is wrong.** FP check 4 claims two independent signals. Strip the self-citation and there is exactly one: the Tor exit node. The correlation analysis independently confirms this -- all four chain legs returned "Proposed, not found."
- **The check designed to catch this is what it defeated.** FP check 4 asks whether the alert rests on independent evidence "rather than resting on a single uncorroborated detection." The honest answer here is that it rests on one detection plus one enrichment. The check passes by counting the report as its own witness.
- **It escalated a real action.** Revoke Session is gated at 0.75 and cleared at 0.90. Some part of that 0.90 is the self-citation. A live AWS session revocation is being recommended on a confidence figure partly composed of the recommendation itself.
- **It is invisible at volume.** Nobody reading one report notices that two hex strings match. At a thousand incidents a day this produces a verdict history that appears densely corroborated and contains no independent facts at all.

This is a same-week fix and it was already on the list.

Verdict records need a provenance class (`HUMAN_CONFIRMED / AI_UNREVIEWED`), and `get_identity_verdict_history` needs to exclude the current incident ID from its own result set. The second of those is a one-line filter. Whatever else is on the roadmap, ship that filter first.

03 What is still broken, and what is new

UNRESOLVED ACROSS ALL FOUR DOCUMENTS

ID	DEFECT	AGE	NOTE
D-01	MITRE technique T1685.002 does not exist	4 docs	Now in the narrative prose, not just the mapping block. T1078.004 elsewhere in the same document is valid, so the capability exists.
D-10	Qualifiers q1 and q2 still read "Not confirmed"	4 docs	Flagrant now: q1 asks whether a replacement trail was created, and the report answers it twice -- a CreateTrail search and a lifecycle check. State is still not written back.
D-14	No raw records; timeline holds one entry	4 docs	With a lifecycle state, an IP verdict and a cited prior record available, there is more to plot than ever. Still one line. "6 records" still unspecified.
--	A tool is called and its output never used	3 docs	get_aws_activity appears in Evidence Gathered with no parameters and no result anywhere in the narrative. Same pattern as get_identity_baseline in v3 and get_detection_dependencies in v1.
--	No identity-provider telemetry	4 docs	The report concludes "compromised principal" and recommends investigating the root cause of the compromise, with zero authentication data.

NEW IN THIS BUILD

ID	DEFECT	SEV
M-01	Self-citation as corroborating evidence. See page 3.	CRIT
M-02	Two different confidence values for one hypothesis. The banner, the verdict line and the actions block all say 90%. The hypothesis block says H1 - confidence 95%, with H2 at 5% so the pair sums to 100. A reader cannot tell which number the system actually holds. This is the third release in which a single-source-of-truth failure has appeared -- corroboration in v1, actions in v3, the confidence value itself now. The contradiction class keeps migrating rather than dying.	CRIT
M-03	The Overview overstates the rejection. "The activity originating from a Tor exit node [...] effectively rejected the alternative hypothesis." H2's own text raises a fair counter -- that the Tor hit could reflect an overly broad policy -- and nothing in the investigation tests it. Tor justifies strong suspicion, not "effectively rejected" at 95/5 with no authentication data.	HIGH
M-04	Correlation windows still exceed the searches. Legs are scoped at 2h and 4h; the widest search executed reached 3600 seconds. Carried over unchanged from the previous assessment.	HIGH

04 Calibration: real machinery, unreliable number

Calibration was the metric that went backwards last time. It has recovered sixteen points, and the reason is worth being precise about, because what improved is not what most teams would have improved.

What got fixed: confidence is now wired to consequences.

The Recommended Actions block states thresholds and shows its working -- Revoke Session is gated at 0.75, the raw 0.35 match did not clear it, the Tier 3 result at 0.90 did, and human approval is still required. Each action also states why it is warranted at its own tier. That is a confidence-to-action policy matrix, which is precisely what a confidence number is *for*. Before this build the figure was decorative; it now does work.

	v1	v3	v4	ASSESSMENT
Rule match confidence	0.35	0.35	0.35	Unchanged across all three.
Verdict confidence	0.80	0.95	0.90 / 0.95	Stated at two values in one document. See M-02.
Hypotheses considered	2	2	2	Still a forced binary in every release.
Confidences sum to	100	100	100	Still asserting exhaustiveness over a non-exhaustive pair.
Nulls counted as support	3 of 3	2 of 5	2 of 3 (H2)	Improving, but H2's rejection still leans on absence.
Action at this confidence	None	None	3 gated	The structural fix of this release.

Why 38 and not 60.

The machinery is right and the number running through it is not. It is stated twice at two values, it is partly composed of the report citing itself, and it still comes from a two-hypothesis set whose confidences are assigned rather than computed. **A correct policy matrix fed an unreliable number produces confidently wrong automation instead of confidently wrong prose** -- which is a more consequential failure mode, not a less one, because it now terminates in a session revocation rather than a paragraph.

THE SENSITIVITY THAT MATTERS

Rescoring with the self-citation discounted -- corroboration counted as one signal rather than two, which is what the evidence actually supports -- moves the composite from **56.2** to **53.0**. That 3.2-point spread is the price of one unfiltered database query. It is also the cheapest points on the entire roadmap: a filter excluding the current incident ID from its own verdict-history lookup is a one-line change, and it is worth more than any remaining model work.

05 Scorecard across three builds

METRIC	v1	v3	v4		vs v3	BAND	NOTE
Reasoning	42	55	62		+7	SUPERVISED	Tor is a real discriminator and the lifecycle check is decisive. "Effectively rejected" overstates, and two of three H2 disconfirming items are still nulls.
Investigation	31	48	66		+18	SUPERVISED	Current state confirmed, deletion resolved, trail and region named, IP enriched. The three biggest gaps across two assessments closed in one release.
Correlation	27	62	60		-2	SUPERVISED	All four legs honestly "Proposed, not found"; the Disposition admits a single alert. Barely exercised, and 2h/4h legs still rest on 1h searches.
Hypothesis management	38	48	52		+4	UNRELIABLE	H2 contains a genuine steelman (Tor as false positive or overly broad policy). Still two hypotheses, still 95/5, no legitimate-user-on-Tor branch.
False-positive analysis	60	68	66		-2	SUPERVISED	"Not applicable" with a stated coverage reason is three-valued logic finally appearing. Offset because check 4 passes on two signals, one being this document.
Explainability	52	58	64		+6	SUPERVISED	Trail, region, IP and lifecycle state are concrete and the action thresholds carry rationale. Undercut by two confidence values for one hypothesis.
Detection engineering	36	42	52		+10	UNRELIABLE	Suggestion 1 (Tor plus StopLogging) is specific and good. Suggestion 3 correctly uses T1078.004, a real ID. Suggestion 2 still proposes a 5-10 minute window.
Recursive investigation	48	56	58		+2	UNRELIABLE	Good multi-hop chain: anchor, lifecycle, IP reputation, verdict history. It retrieved its own record on the last hop and did not notice.
Evidence quality	26	44	52		+8	UNRELIABLE	Lifecycle status and IP reputation are strong, verifiable artifacts. get_aws_activity was called with no visible parameters and no visible output.
Confidence calibration	28	22	38		+16	FAILING	Action thresholds are explicit, wired and defensible. The number itself is stated twice at two values and is inflated by self-citation.
Cross-platform investigation	18	26	40		+14	FAILING	lookup_ip is genuine external threat intel and same_asset_only is False. Still no IdP query while concluding the principal is compromised.
Response automation	12	38	68		+30	SUPERVISED	Three actions with thresholds, rationale and human-approval gating. A real policy matrix, and the best single fix across three releases.
Data integrity	23	48	47		-1	UNRELIABLE	The fabricated technique name is gone and ARNs are consistent. Offset by a fourth T1685.002, a 90-versus-95 mismatch and a self-referential citation.
Production readiness	33	45	56		+11	UNRELIABLE	Concrete artifacts, gated actions, honest coverage gaps. Still blocked by an invented identifier and evidence that cites the report itself.
WEIGHTED COMPOSITE	33.2	47.7	56.2		+8.6	UNRELIABLE	

Grey = v3 level, colour = movement. Marker at 60 is the trust threshold. **Six metrics now cross it; in v3 two did, in v1 none.** Framework 61.0, substance 53.0 -- ratio **1.15**, from 2.04 at first assessment.

06 How much of this is capability, and how much is the incident?

This is a different incident from the previous two, in a different tenant, with a different principal. Comparing scores across it requires care, and the honest read is that most but not all of the gain is real.

GAIN	ATTRIBUTION	REASONING
Response automation +30	Capability	A policy matrix with thresholds is infrastructure. It would have fired identically on the previous incident.
Investigation +18	Capability	get_trail_lifecycle_status is a new tool that answers questions the previous incident also asked and could not answer.
Cross-platform +14	Mostly capability	lookup_ip is a real new integration. The previous report also had an IP available and did not enrich it.
Calibration +16	Capability	Thresholds and gating are structural, not scenario-dependent.
Reasoning +7	Mostly the incident	A Tor exit node is a far easier discriminator than the ambiguous administrative-action case in the previous two reports. The reasoning did not have to work as hard.
Correlation -2	Untested	Nothing correlated this time, so the capability was barely exercised. The honest "Proposed, not found" labels are good behaviour but demonstrate less than v3 did.
Detection eng +10	Mixed	Suggestion 1 is Tor-specific and only exists because this incident had Tor in it. Suggestion 3 using a valid technique ID is genuine capability.

Roughly six of the 8.6 points are capability. Two to three are the incident being easier.

That is a good ratio and it should not be talked down -- but it means the true comparable position is nearer **53-54** than 56.2, and it means **you have not yet re-tested the hard case**. The genuinely useful next measurement is re-running the original ambiguous incident, where there is no Tor node and no obvious discriminator, through this build. If it still reaches 0.95 on absence-of-evidence there, the epistemics have not moved; the evidence just got easier.

WHAT TO DO NEXT

WHEN	ACTION	WHY
Today	Filter the current incident ID out of get_identity_verdict_history results.	One line. Removes the most serious defect in this document and 3.2 points of inflated score.
This week	Verdict provenance classes; AI_UNREVIEWED may not satisfy an independence check.	The general form of the same bug. Without it the filter is a patch, not a fix.
This week	Single-source the confidence value. One number, one place, rendered everywhere.	Third release with a single-source-of-truth failure. Fix the pattern, not the instance.
This week	ATT&CK; enum validation. Map T1685.002 to T1562.008.	Fourth document. It is days of work and it is the first thing an evaluator checks.
Next sprint	Write qualifier answers back to q1/q2 state.	You already do the work in every release and never record it.
Next sprint	Match search windows to the correlation legs you report against.	Your honest labels are only as good as the searches beneath them.
Next sprint	Re-run the original ambiguous incident through this build.	The only way to separate capability from scenario difficulty.

07 Verdict

Three assessments ago this product scored 33.2 and reached a probably-correct conclusion through absence-of-evidence reasoning, with a corrupted account ID in its evidence and no remediation advice at all. It now scores 56.2, names the trail and region, confirms the trail is still disabled, enriches the source IP, reports its own coverage gaps honestly, and recommends three threshold-gated actions with human approval. That is a real engineering trajectory and the team executing it is clearly reading the findings and acting on them.

It is also a report whose false-positive gate was cleared, in part, by the report itself. The single most important structural warning from the last assessment -- that improving evidence retrieval without adding evidence provenance would turn a projected risk into a live one -- was not acted on, and the outcome is a document citing its own incident ID as prior corroborating evidence in support of revoking a production session. **The pattern across all three builds is consistent: what gets built gets built well, and what gets deferred gets worse.**

The framework-to-substance ratio closing from 2.04 to 1.15 is the number to be pleased about. The presentation layer and the evidence layer have converged, which was the central criticism of the first assessment and is now substantially answered. What remains is narrower than what has been fixed: an identifier that has never been validated, a confidence value that exists in two places at two values, and a verdict store with no notion of where its verdicts came from. None of those require a model change, a new integration, or a research result.

56.2 / 100 -- top of UNRELIABLE. Four points from the trust threshold, and one query filter from earning them honestly.

The three P0 items on page 7 are days of work in total. They would put this build at roughly **60** -- across the boundary where an analyst verifies the work rather than redoing it, which remains the only threshold that changes what you can honestly sell. Do not ship the word "autonomous" while a session revocation can be justified by a self-citation.

LIMITATIONS OF THIS ASSESSMENT

- **Different incident, so the delta is not clean.** v1 and v3 were the same incident; this is a new one with an easier discriminator. Page 7 attributes the gain explicitly, but the comparison is weaker than the previous one and should not be read as a controlled measurement.
- **Still synthetic.** The tenant is named xplat-synth-test. Sparse context stores and clean telemetry continue to flatter the product; a populated customer estate with real change tickets will behave differently, and that is where the absence-of-evidence habit will be tested properly.
- **Exported PDF only.** No access to the verdict store, so it cannot be confirmed whether KB_RECORD_ID 15369769f5175238 is a genuine self-reference or a rendering error that substituted the current ID for a real prior record. Both are serious; the second is cheaper to fix. The report gives a reader no way to distinguish them, which is itself the finding.
- **MITRE checked against ATT&CK Enterprise.** T1685.002: still no entry. T1562.008 remains correct. T1078.004, cited in Detection Opportunity 3, is valid and correctly applied.